


Hierarchical attacker-defender policy

A high-level decision identifies where to act; a low-level allocator determines the bounded action intensity.



Training reward measures optimization progress; it does not by itself establish Nash equilibrium.