


One ODE-based environment transition

At epoch t_k , the policy acts once; the simulator then advances the continuous state to t_{k+1} .



Ordering is part of the model: reset-before-flow and flow-before-reset generally produce different transitions. Running and impulse costs are recorded separately before they are combined in the reward.